

# Battle of Time Series: SQL Server 2022 vs. ADX vs. InfluxDB vs. TimescaleDB

**DATA**  
SATURDAYS



# Sponsors



# Gianluca Sartori



Founder of



@spaghettidba

 spaghettidba.com

 quantumdatis.com

Independent SQL Server  
consultant

Data platform MVP

Works with SQL Server since 7.0

DBA @ Scuderia Ferrari

1. What is Time Series data?
2. Options on the market
3. The contestants
4. The challenges
5. The winner

# What is Time Series data?

## What is **time series data**?

*“Time series data (or time-stamped data) is a collection of observations for a single object (entity) at different time intervals.”*


## What is a **time series database**?

*“A time series database (TSDB) is a database optimized for **time series data** and for measuring change over time.”*

Time series data is obtained by performing repeated measurements over time

## Examples:

- Atmospheric temperature
- Stock prices
- CPU usage %
- Emails/sec
- Sensor data



## How is this different from relational data?

- Continuous Stream of data in time order
- Bulk uploads of large sets of data
- High volume data
- Data is append-only – no updates
- Delete large volumes of data when it goes out of scope
- Downsampling and aggregating high resolution data to save space



## Why does it matter?

### Aircraft Engine

10 terabytes every 30 min.  
14 million hours of flight in  
2018  
2 engines  
= 560 million TB

**2020s**

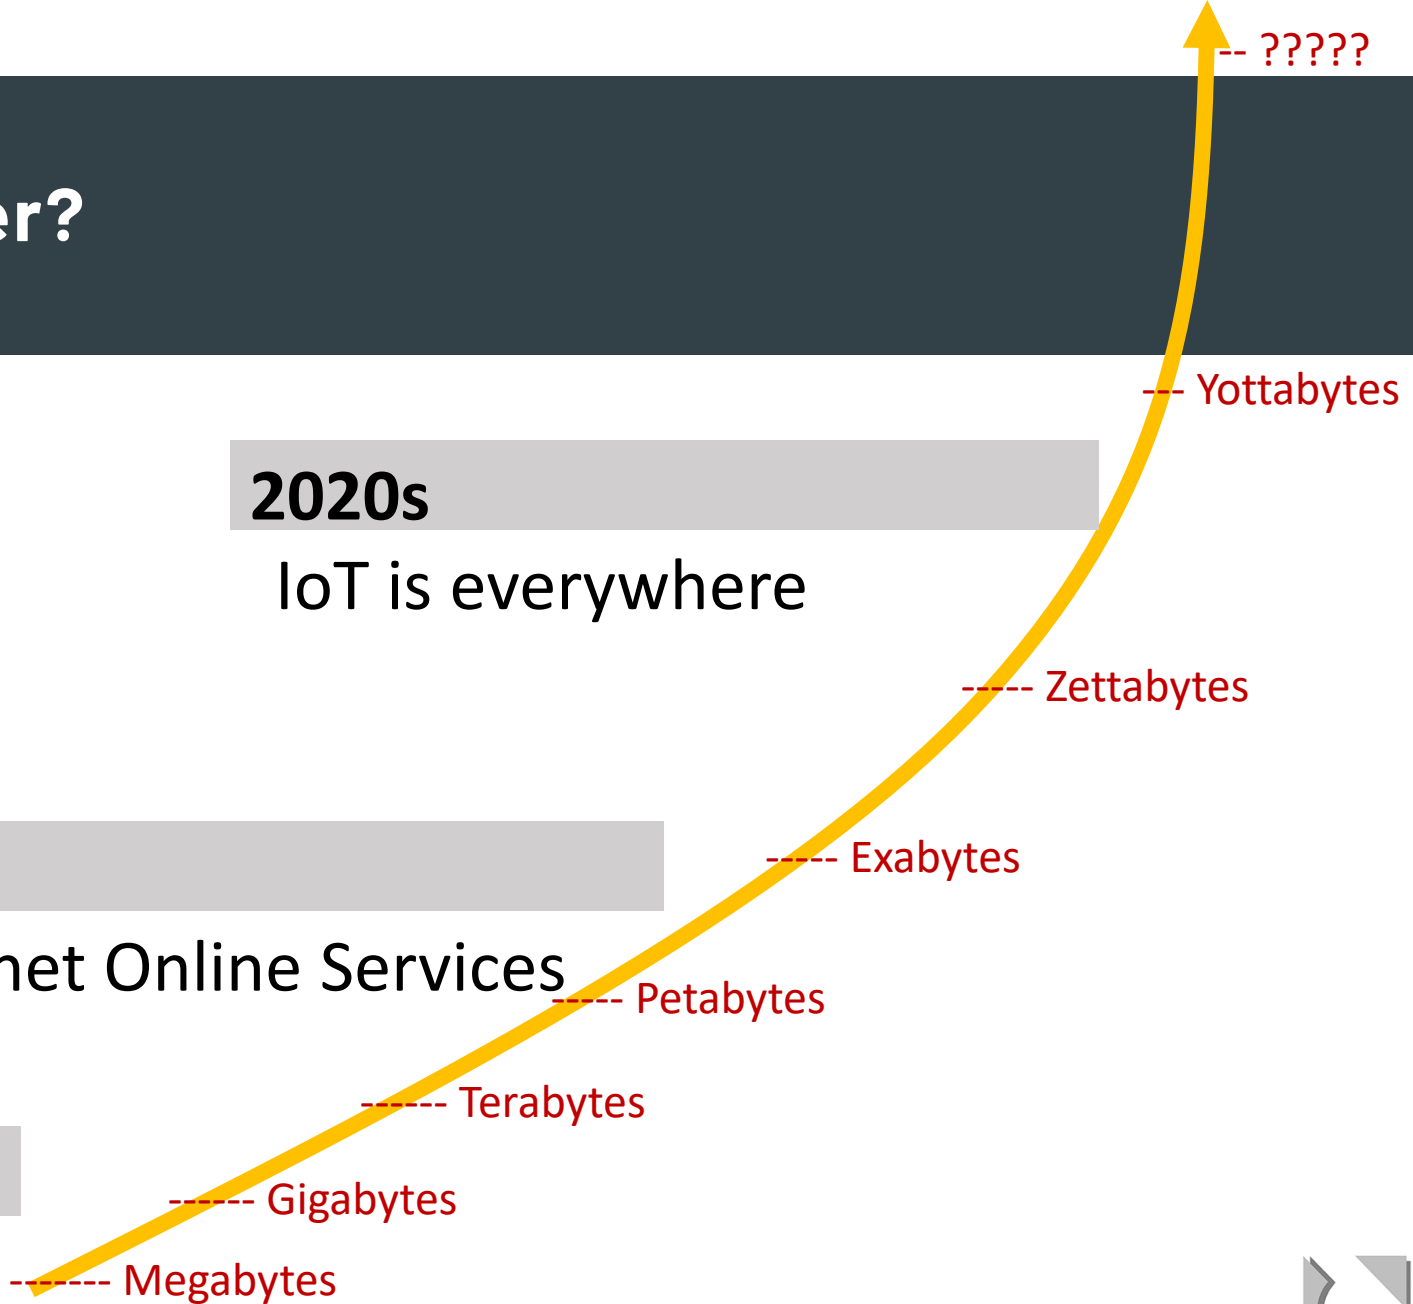
IoT is everywhere

**2000s**

Internet Online Services

**1970s**

Early Relational Databases



# Time Series Database Options

RANK	DBMS	SCORE	
NOV 2024		NOV 2024	24 MONTH
1	InfluxDB	21.47	-8.55
2	KDB+	7.06	-1.72
3	Prometheus	6.92	+0.61
4	Graphite	4.91	-0.45
5	TimescaleDB	3.68	-0.88
6	QuestDB	2.91	+2.01
7	Apache Druid	2.70	-0.48
8	DolphinDB	2.60	+0.98
9	TDEngine	2.27	+1.08
?????	Azure Data Explorer	????	?????




# Battle of TimeSeries!!!!



**Timescale**



Azure Data Explorer



Microsoft®  
**SQL Server®**



***influxdb***

- All-round, enterprise-ready relational database
- Not a time-series database
- Has some of the ingredients
- Has a specialized edition for time-series data called Azure SQL Edge



- Open Source Time-Series database
- Built on top of PostgreSQL as an extension
- Optimized to ingest, store and query time-series data
- Stores data in «HyperTables», special tables composed of «chunks»
- Time buckets partition data by timestamp
- Allows compression for
- Can also store relational data and all types of data supported by PostgreSQL







- Fully managed, high performance, big data analytics platform
  - Can store and analyze structured, semi-structured and unstructured data
  - Uses a «relational» model (tables, columns and rows) with fixed strongly-typed schemas
  - Tables are stored in databases
  - Clusters contain databases
  - Log analytics, IoT, time series data
- ➔ not strictly a time-series database!





Most popular time series database

Designed and optimized for time series data

Uses optimized TSM storage engine with columnar compression


Built-in retention policies


Open Source with permissive licensing (MIT)

Written in GO

Cross platform:

windows, linux, macOS, docker,  
kubernetes, raspberryPI





1. Licensing
2. Installation
3. Tooling
4. Storage Efficiency
5. Retention Policies and Downsampling
6. Query Capabilities

The rules:

- 0 to 3 points awarded
- I am the sole judge :)



# Licensing





## SQL Server

- Commercial Product – public pricing
- Closed Source
- Has cloud option(s) – public pricing
- Has free editions with limitations

2



## InfluxDB

- Commercial Product – secret pricing
- Open Source
- Has cloud option – public pricing
- Has extremely limited community edition
- Different versions with different limitations

0



## TimescaleDB

- Commercial Product – secret pricing
- Open Source
- Has cloud option – public pricing
- Has free editions with no limitations

3




## Azure Data Explorer

- Commercial Product – public pricing
- Closed Source
- Only available in the cloud
- Offers a limited free tier

1

# Installation





## SQL Server

- Cross-platform
- Works in containers
- GUI setup – Unattended setup
- Extensive and accurate documentation

3



## InfluxDB

- Cross-platform
- Works in containers
- No setup required – single executable
- Nearly non-existent documentation

2



## TimescaleDB

- Cross-platform
- Works in containers
- GUI setup – Unattended setup
- Extensive and accurate documentation *with OSS factor*

1



## Azure Data Explorer

- No installation required
- Extensive and accurate documentation

1



# Tooling





## SQL Server

- SSMS
- Azure Data Studio
- Code extension
- Countless alternatives (some OSS)

3



## InfluxDB

- Web UI - Chronograf
- CLI (influx.exe)
- Unofficial – unsupported OSS tools

0



## TimescaleDB

- pg\_admin
- Code extension
- Countless alternatives (some OSS)

2



## Azure Data Explorer

- Web UI - [dataexplorer.azure.com](https://dataexplorer.azure.com)
- Kusto Explorer
- Azure Data Studio

1

# Storage Efficiency



Big Heaps or Clustered Indexes are problematic

- Huge in size
- Hard to query
- Hard to maintain

Alternatives?

- Data compression
- Clustered Columnstore Indexes



- Each column is stored independently
- Compresses data very efficiently
- Suitable when lots of repeating data exists
- Good for bulk operations
- Don't play well with UPDATES and DELETES

## DEMO!

SQL Server CPU usage percent from `sys.dm_os_ring_buffers`

time	SQL Server CPU	Other CPU
08:00	10	3
08:01	15	2
08:02	18	6
08:03	73	1

# InfluxDB Series – An Example

This is saved in a **measurement** (let's call it **CPUHistory**) with **two series**:

time	SQL Server CPU	Other CPU
08:00	10	3
08:01	15	2
08:02	18	6
08:03	73	1

[CPUHistory, SQL Server CPU]

time	SQL Server CPU
08:00	10
08:01	15
08:02	18
08:03	73

[CPUHistory, Other CPU]

time	Other CPU
08:00	3
08:01	2
08:02	6
08:03	1

# Series – Tags

If we add **tags**, we will get **one series for each tag value**:

time	ServerName	SQL Server CPU	Other CPU
08:00	ACCOUNTING	10	3
08:00	CRM	37	7
08:01	ACCOUNTING	15	2
08:01	CRM	48	5
08:02	ACCOUNTING	18	6
08:02	CRM	41	9
08:03	ACCOUNTING	73	1
08:03	CRM	28	7

[CPUHistory, **CRM**, OtherCPU]

time	Other CPU
08:00	7
08:01	5
08:02	9
08:03	7

[CPUHistory, **ACCOUNTING**, SQLServerCPU]

time	SQL Server CPU
08:00	10
08:01	15
08:02	18
08:03	73

[CPUHistory, **ACCOUNTING**, OtherCPU]

time	Other CPU
08:00	3
08:01	2
08:02	6
08:03	1

[CPUHistory, **CRM**, SQLServerCPU]

time	SQL Server CPU
08:00	37
08:01	48
08:02	41
08:03	28

Series are stored independently

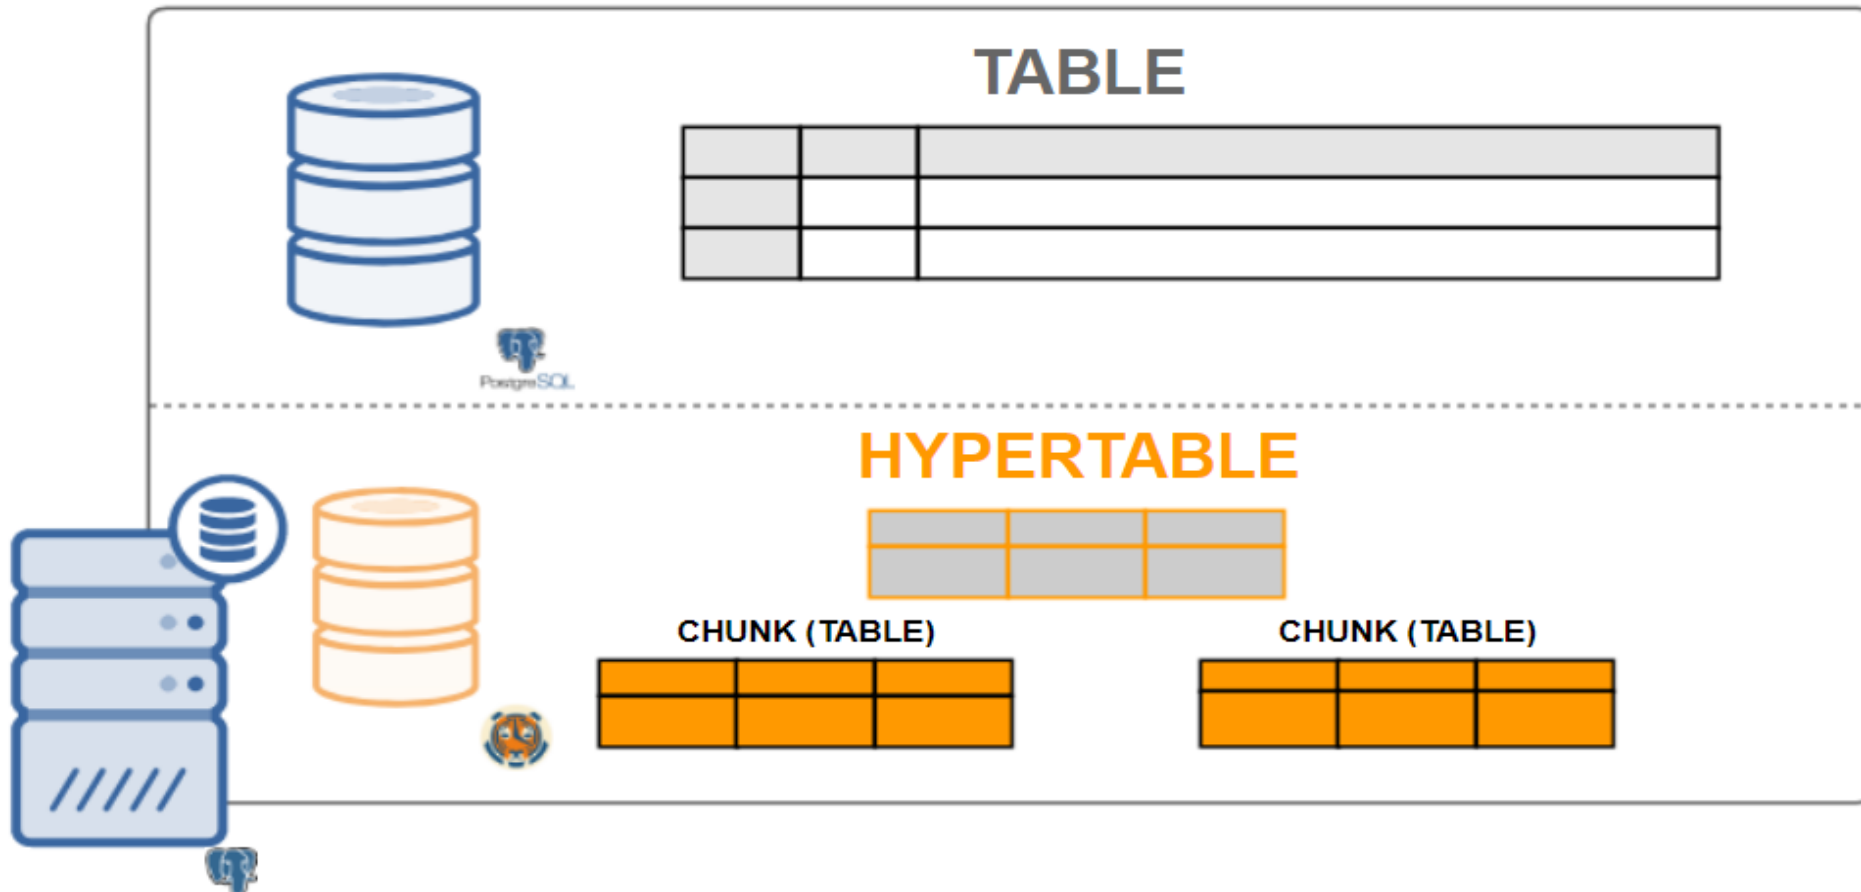
We get one series for each combination of tag values → **cardinality**

v1 & v2 Storage engine is called **TSM** (Time series Structured Merge tree)  
derived from **LSM** (Log-Structured Merge tree) → **Cassandra**

v3 Storage is based on **Parquet**

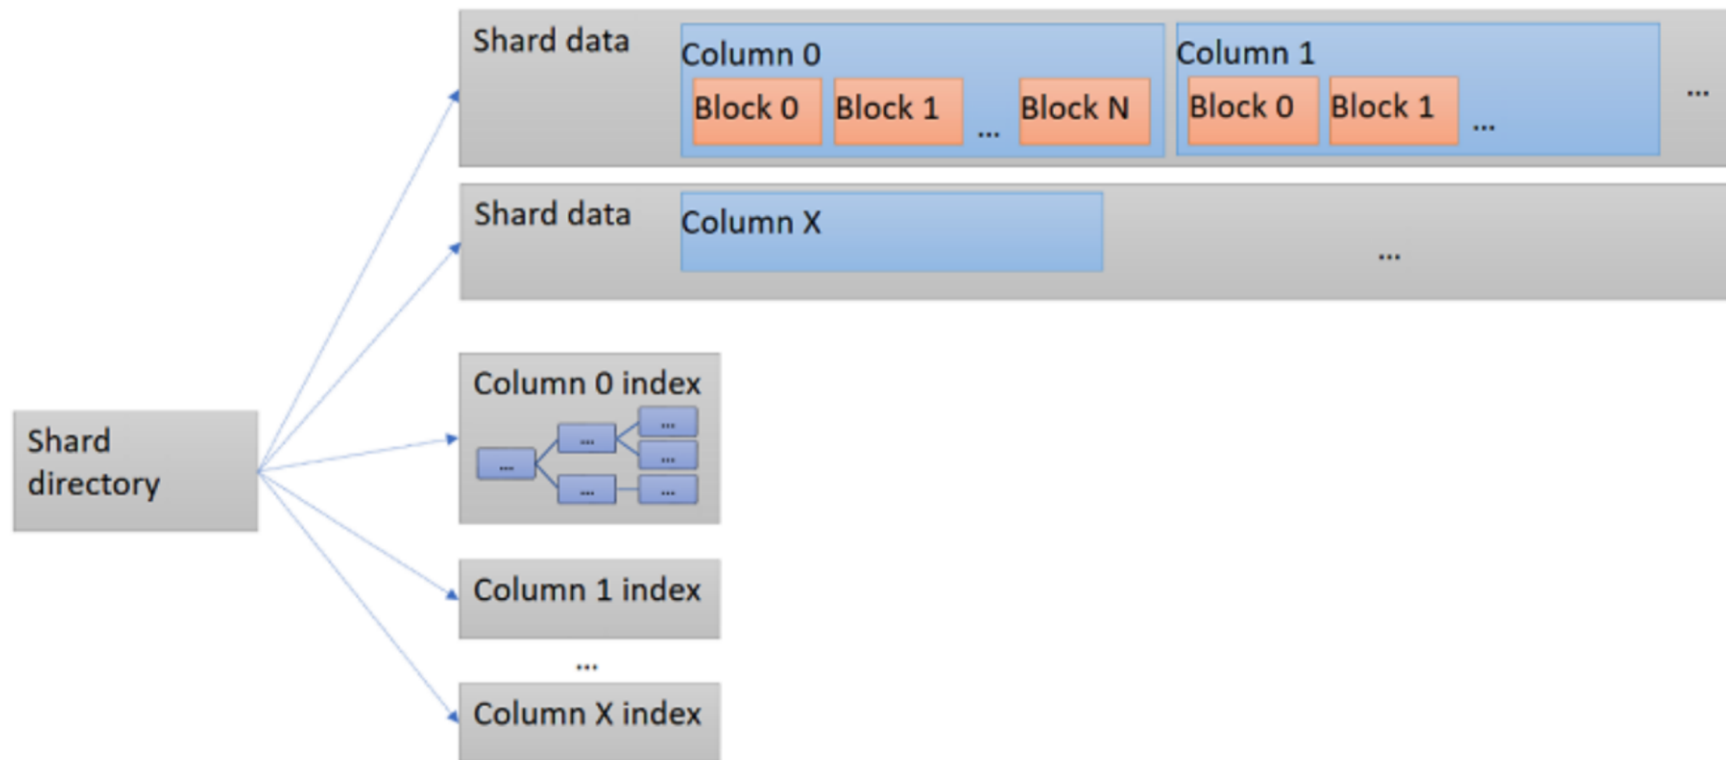
Available in the InfluxDB cloud  
(Soon?) available on-prem

# TimescaleDB – Storage Engine




- Separates storage and compute resources
- Persistent data resides in Azure Blob Storage
  - Data is stored in **extents** (shards)
  - Extents are spread across cluster nodes
  - Extents are cached in SSD and memory
  - Data is compressed with columnar compression
- Compute uses a cache for persistent storage
- Row store is used when ingesting data (streaming ingestion)

Compressed column store with free text support and full-text inverted index





Data size for SQL Server performance counters, 2 instances,  
15 days of data, 82 million rows





## SQL Server

- Terrible with regular tables
- Better with clustered columnstore
- Partitioning affects efficiency

1



## InfluxDB

- Stellar compression ratios
- Unaffected by sharding

3



## TimescaleDB

- Terrible with regular tables
- Regular hypertables are just as inefficient
- Compressed hypertables are fine

2







## Azure Data Explorer

- Uses an undocumented columnar compression algorithm
- Stores data efficiently

2

# Retention Policies and Downsampling



	9		8		11		6
---	---	---	---	---	----	---	---



## SQL Server

- Data needs to be deleted manually
- Partitioning makes it nearly-instantaneous
- Requires scheduling jobs

0



## InfluxDB

- Retention policies are built-in
- Downsampling can be achieved with continuous queries in v1 and with jobs in v2 and v3

3



## TimescaleDB

- Retention policies are built-in
- Downsampling is implemented with materialized views

3



## Azure Data Explorer

- Retention policies are built-in
- Data gets soft-deleted when expired
- Downsampling is implemented with materialized views or with Power Automate
- Not really a time-series database

1

# Query Capabilities





InfluxDB supports 2 query languages:

**InfluxQL** → similar to SQL, born with version 1, still available in v2

```
SELECT MEAN(SQLServerCPU) AS avg_cpu  
FROM CPUHistory  
WHERE time > now() -1d  
GROUP BY server_name
```

**Flux** → doesn't look like SQL (more like KQL), default in v2, also available in v1

Supports «advanced» query features like joins...

```
from(bucket:"demo/autogen" )
  | > range(start: -1d)
  | > filter(fn: (r) => r._measurement == "CPUHistory")
  | > filter(fn: (r) => r._field == "SQLServerCPU")
  | > mean()
  | > group(columns: ["server_name"])
  | > yield(name: "avg_cpu")
```

## InfluxDB v3 supports SQL

InfluxDB Cloud Serverless uses the **Apache Arrow DataFusion** implementation of SQL

Supports JOINS

Supports CTEs

Supports windowing functions

... Supports everything you would expect from a decent database



InfluxQL allows grouping by time:

```
SELECT MEAN(value) AS avg_lazy_writes_sec  
FROM "sqlserver_performance"  
WHERE counter = 'Lazy writes/sec'  
      AND sql_instance = 'SQLCSRV04:SQL2017'  
      AND time > now() - 1d  
GROUP BY time(1h)
```

... GROUP BY time(1h)

time	lazy_writes_sec		time	avg_lazy_writes_sec
16-10-2022 10:00:00	398091	10:00	16-10-2022 10:00:00	398,091.00
...	...		16-10-2022 11:00:00	400,187.50
16-10-2022 10:59:59	395644	11:00	16-10-2022 12:00:00	406,095.00
16-10-2022 11:00:00	432211		16-10-2022 13:00:00	411,167.76
...	...	12:00	16-10-2022 14:00:00	417,963.40
16-10-2022 11:59:59	452325		16-10-2022 15:00:00	425,683.27
16-10-2022 12:00:00	456541	13:00	16-10-2022 16:00:00	433,001.40
...	...		16-10-2022 17:00:00	441,937.02
16-10-2022 12:59:59	433505	...	16-10-2022 18:00:00	449,937.73
16-10-2022 13:00:00	456654		16-10-2022 19:00:00	450,737.81
...	...		16-10-2022 20:00:00	453,467.40
16-10-2022 13:59:59	456788		16-10-2022 21:00:00	455,703.76
16-10-2022 14:00:00	417545		16-10-2022 22:00:00	457,270.19
...	...		16-10-2022 23:00:00	458,743.47
			17-10-2022 00:00:00	461,533.64
			...	...

InfluxQL allows filling intervals with data:

```
SELECT MEAN(value) AS avg_lazy_writes_sec  
FROM "sqlserver_performance"  
WHERE counter = 'Lazy writes/sec'  
      AND sql_instance = 'SQLCSRV04:SQL2017'  
      AND time > now() - 1d  
GROUP BY time(1h) FILL(NULL)
```

... FILL(NULL | VALUE)

time	value
16-10-2022 10:35:45	398091
16-10-2022 10:36:01	398091
16-10-2022 12:16:15	406095
16-10-2022 12:16:30	406095
16-10-2022 13:06:45	411167
16-10-2022 13:07:01	411167

No data  
for  
11:00



time	avg_lazy_writes_sec
16-10-2022 10:00:00	398,091.00
16-10-2022 11:00:00	NULL
16-10-2022 12:00:00	406,095.00
16-10-2022 13:00:00	411,167.76

The new DATE\_BUCKET function in SQL Server 2022 helps implement GROUP BY TIME

```
SELECT DATE_BUCKET(hour, 1, time) AS time, AVG(value) AS avg_lazy_writes_sec
FROM [dbo].[sqlserver_performance_cci_partitioned]
WHERE time >= '2022-10-01'
      AND time < '2022-10-02'
      AND counter = 'Lazy writes/sec'
      AND sql_instance = 'SQLCSRV04:SQL2017'
GROUP BY DATE_BUCKET(hour, 1, time)
ORDER BY 1
```

The new GENERATE\_SERIES function in SQL Server 2022 helps implement FILL

```
SELECT DATEADD(hour, value, '2022-10-01') AS time  
FROM GENERATE_SERIES(0, DATEDIFF(hour, '2022-10-01', '2022-10-02') - 1)
```

# DEMO!

- ADX supports the Kusto Query Language (KQL)
- It's a brand new language, no similarities with SQL
- SQL is not supported but you can convert SQL to KQL with EXPLAIN
- KQL is simple and powerful
- ... but it's not SQL!!!
- SQL support for KQL databases [is coming in 2025!!!](#)

```
1 Perf
2 | where TimeGenerated >= ago(10m)
3 | where CounterName == "% Free Space"
4 | project PerfComputer = Computer
5         , CounterName
6         , CounterValue
7         , PerfTime=TimeGenerated
8 | join ( InsightsMetrics
9         | where TimeGenerated >= ago(10m)
10        | project IMComputer = Computer
11            , Namespace
12            , Name
13            , Val
14            , IMTime=TimeGenerated
15        )
16 on $left.PerfComputer == $right.IMComputer
17
18
```

## DEMO!



# TimescaleDB – Query capabilities

PostgreSQL has a rich SQL dialect called PL/pgSQL (or just pgSQL) which handles most queries with ease.

TimescaleDB adds **hyperfunctions** to handle common time-series use cases

GROUP BY TIME is achieved with **time\_bucket()**

FILL is achieved with **time\_bucket\_gapfill()**

Carry over of the last value is implemented with **locf()**

# DEMO!

	10		11		14		8
---	----	---	----	---	----	---	---



## SQL Server

- Supports T-SQL
- Has gained (some) time-series capabilities with SQL Server 2022

1



## InfluxDB

- Has specific time-series query features
- InfluxQL
- Flux
- SQL

3



## TimescaleDB

- Supports pgSQL
- Has specific time-series query features

3



## Azure Data Explorer

- Has specific time-series query features
- Supports KQL
- Has a SQL endpoint

2

	10		11		14		8
---	----	---	----	---	----	---	---

**TimescaleDB!**



Questions? Ask me!

[spaghettidba@sqlconsulting.it](mailto:spaghettidba@sqlconsulting.it)

Get slides and demos from GitHub:

<https://github.com/spaghettidba/CodeSamples>