

Three Things you don't want
to store in a database
(but sometimes have to)

Gianluca Sartori

Founder at Quantumdatis



Data Platform MVP



Works with SQL Server since version 7

DBA @ Scuderia Ferrari



Blog: spaghettidba.com

Twitter: [@spaghettidba](https://twitter.com/spaghettidba)



Agenda

- What are the rules?
- Logs
- Pictures
- Dynamic Attributes

A large blue triangle is positioned on the left side of the slide, pointing towards the bottom right corner.

The problem with Normal Forms

The key, the whole key, nothing but the key

- 1NF:
A primary key, atomic attributes only
- 2NF:
Every attribute depends on the whole key
- 3NF:
Every attribute depends only on the key

A large blue triangle pointing downwards from the top-left corner, covering the left side of the slide.

Bad idea #1
Storing logs in the database

What do you need to do with logs?

- Write efficiently
 - Writing doesn't have to affect the application performance
 - Written synchronously
 - Mostly written and never read
- Query & analyze
 - You read logs for troubleshooting
 - Usually read sequentially
 - Analysis
- Apply retention policies
 - Keep for a reasonable time, then delete

Logs don't belong in a database

- It's not really 1NF compliant
 - You don't really care about the whole message, but parts of it
 - Huge queries WHERE message LIKE '%something%'
- It's not transactional data
 - You don't need logs to be ACID compliant
 - When you delete old logs the database log will explode
 - You need to log when the database is not reachable
- There is no benefit
 - RDBMSs deal with relational data very well
 - Logs are... logs. What can a RDMBS offer?

What are the alternatives?

- Files
 - Easy to write to
 - Performance is great
 - Available when the database is offline / unreachable
 - Easy to read
 - Text editors & low tech
 - Human readable
 - Easy to query
 - GREP
 - [LogParser](#)
 - Easy to delete
 - Just delete old files

What are the alternatives?

- Specialized logging databases

- Logstash
- Loki
- Loggly
- ElasticSearch
- InfluxDB
- GrayLog
- DataDog
- Splunk

- Strengths:

- Optimized for time
- Analysis



elasticsearch



Logstash

Why in a RDBMs?

- Easy - lazy
 - Not different from any other type of data
 - No need to learn new techniques / technologies
- Convenient for filtering & displaying data
 - Filter with SQL queries
- Can leverage RDBMS features
 - High Availability
 - Replication
 - Backups

Challenges

- High cardinality
 - Storage
 - Querying
 - Retention

Possible solutions


- Choosing the right storage
 - Heaps?
 - Clustered Indexes?
 - Compression?
 - Columnstore indexes?
 - Partitioning?

DEMO


LOGS

Heaps vs Clustered Indexes

Space used




Time taken




Compressed vs Uncompressed

Space used




Time taken

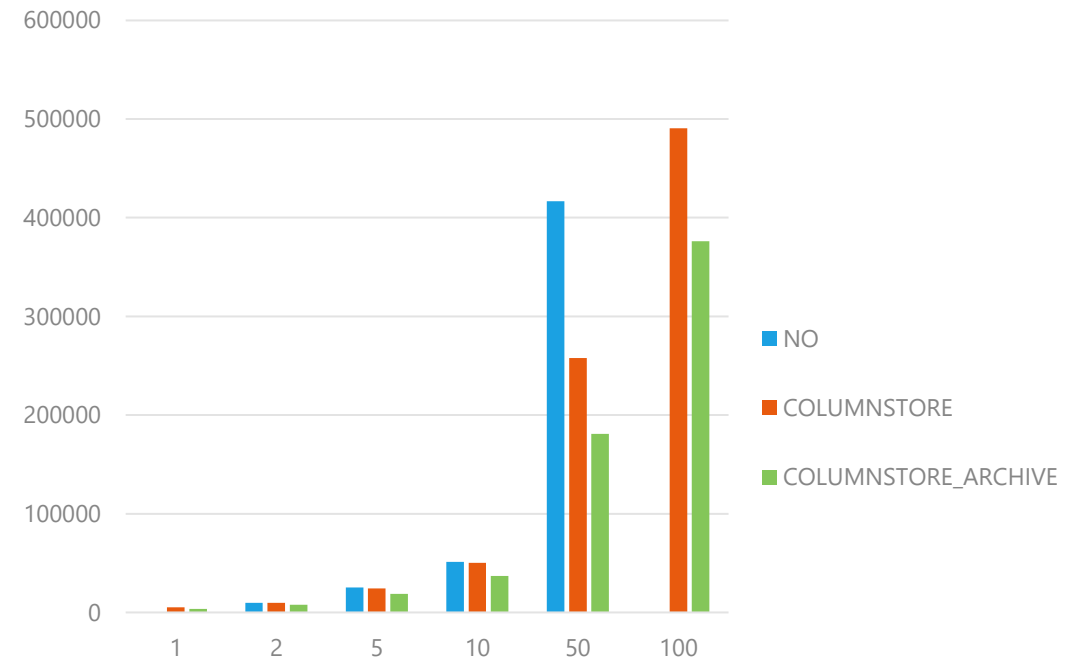


Columnstore vs Rowstore

Space used



Time taken



Wrap up

- Columnstore wins because of the shape of the data
- If logs are not repetitive enough it might be a problem
- Careful with LOB data
- Use partitioning to compress rows
- Use partitioning to delete old data
- Don't try to delete rows on compressed segments: columnstore doesn't like deletes

A large blue triangle pointing downwards from the top-left corner of the slide, extending diagonally across the left side.

Bad idea #2

Storing pictures in the database

What do you need to do with pictures?

- It's not only pictures, it's binary files in general. Pictures is a very common example
- Write once, read as a whole
 - Mostly read, seldom written
 - Always read entirely
- Pass to other applications
 - Read the contents, pass to a web server or client app
 - Always read and write entirely

Pictures don't belong in a database

- It's not transactional data
 - You don't need pictures to be ACID compliant
 - When you maintain indexes the database log will explode
- There is no gain, only pain
 - RDBMSs are designed for relational data
 - Pictures don't fit well in pages/extents
 - Throughput is poor

What are the alternatives?

- Files
 - Easy to write to.
 - Performance is great
 - Available when the database is offline / unreachable
 - Easy to visualize/edit
 - Low tech
 - Easy to update/delete
 - Just replace or delete files

What are the alternatives?

- Specialized object storage
 - Azure BLOB storage
 - Amazon S3
 - ... all cloud providers have a solution
 - NetApp
- Strengths:
 - Optimized for analytics
 - History
 - Distributed
 - Scale-out

Microsoft Azure
Blob Storage



Why in a RDBMs?

- Easy - lazy
 - Treat as the rest of the data
 - No need to learn new techniques / technologies
- Can leverage RDBMS features
 - High Availability
 - Replication
 - Backups

Challenges

- Scales very poorly
 - Reads and writes are poor compared to file system
 - Index maintenance is virtually impossible
 - Contributes to backup size

Possible solutions

- Choosing the right storage
 - Heaps?
 - Clustered Indexes?
 - Compression?
 - Columnstore indexes?
 - Filestream

DEMO

PICTURES

Wrap up

- Filestream delivers best performance
- Be careful with missing files (CHECKDB will detect)
- NTFS performance affects filesystem access
 - Disable 8.3 naming
 - Disable indexing
 - Disable lastaccess
 - Dedicated disk
 - Choose correct cluster size
 - Exclusions for antivirus / antimalware
- Skipping TDS might improve performance

A large blue triangle is positioned on the left side of the slide, pointing towards the bottom right corner.

Bad idea #3
Dynamic attributes
in the database

Why dynamic attributes?

- Sometimes it's hard to identify all attributes at design time
- Some applications may have the ability to store dynamic attributes
 - CRM
 - RAD

Dynamic attributes don't belong in a database

- All attributes of entities must be known at design time
- The relational model does requires equal number of attributes for each row
- Some (bad) implementations violate 1NF, 2NF or 3NF

What are the alternatives?

- Document databases
 - Azure Cosmos DB
 - MongoDB
 - Amazon DocumentDB
 - ElasticSearch



Why in a RDBMs?

- Looks «smart»
 - There are many possible implementations (some are very bad)
 - Stays with the rest of the data
- Can leverage RDBMS features
 - High Availability
 - Replication
 - Backups

Possible solutions

- EAV
 - Violates 2NF, 3NF
- XML
 - Violates 1NF
- JSON
 - Violates 1NF
- Sparse columns
 - Lots of NULLs
 - Generic attribute names

DEMO

DYNAMIC ATTRIBUTES

Wrap up

- Dynamic attributes have no place in relational databases
- XML and JSON violate 1NF – square peg, round hole
- EAV is the worst
 - Generic data types → Eg: varchar(4000)
 - No Foreign Keys
 - No CHECK constraints
 - Multiple accesses to the same table
- SPARSE columns are not too bad
 - Doesn't work when every row has completely different attributes

Wrap up

*To store data in a database, you need to know the rules.
In a relational database, the rules are called Normal Forms.*

Ignore them to enter a world of pain.