

Gianluca Sartori



Time Series for Relational People

Gianluca Sartori

Founder at Quantumdatis

Data Platform MVP

Works with SQL Server since version 7

DBA @ Scuderia Ferrari

Blog: spaghettidba.com

Twitter: [@spaghettidba](https://twitter.com/spaghettidba)



Agenda

- What is a Time Series database?
- Options on the market
- InfluxDB
 - Core concepts
 - Installation
 - Working with data in InfluxDB
 - Querying
- Practical uses





What is a
Time Series
database?

Definitions

What is **time series data**?

“Time series data (or time-stamped data) is a collection of observations for a single object (entity) at different time intervals.”

What is a **time series database**?

*“A time series database (TSDB) is a database optimized for **time series data** and for measuring change over time.”*




Time Series Data

Time series data is obtained by performing repeated measurements over time

Examples:


- Atmospheric temperature
- Stock prices
- CPU usage %
- Emails/sec
- Sensor data



Time Series Data


Metrics (Regular)

Measurements gathered at regular time intervals



Events (Irregular)

Measurements gathered at irregular time intervals



Time Series Data

How is this different from relational data?

- Continuous Stream of data in time order
- Bulk uploads of large sets of data
- High volume data
- Data is append-only – no updates
- Delete large volumes of data when it goes out of scope
- Downsampling and aggregating high resolution data to save space



Should I use SQL Server?

Relational databases do not deal very well with time series data

- No built-in storage optimizations
- No easy retention policies enforcement
- No time series specific query capabilities

- **Azure SQL Edge** is designed for time series data

See <https://docs.microsoft.com/en-us/azure/azure-sql-edge/>

- **SQL Server 2022** will have time series capabilities from Azure SQL Edge

See <https://techcommunity.microsoft.com/t5/microsoft-mechanics-blog/what-s-new-in-sql-server-2022/ba-p/2922227>



Time Series Database Options

RANK	DBMS	SCORE		
		APR 2022	24 MOS ▲	12 MOS ▲
1	InfluxDB	30.02	+9.10	+2.85
2	Kdb+	8.78	+3.41	+0.52
3	Prometheus	6.31	+2.00	+0.55
4	Graphite	5.36	+1.90	+0.80
5	TimescaleDB	4.56	+2.79	+1.66
6	ApacheDruid	3.18	+1.27	+0.51
7	RRDtool	2.58	+0.06	+0.12
8	OpenTSDB	1.82	0.00	+0.02
9	DolphinDB	1.62	+1.23	+0.72
10	Fauna	1.42	+0.47	-0.07

Source: DB-Engines

39 Systems in Ranking, April 2022



InfluxDB



Most popular time series database

Designed and optimized for time series data

Uses optimized TSM storage engine with columnar compression


Built-in retention policies

Open Source with permissive licensing (MIT)

Written in GO

Cross platform: *windows, linux, macOS, docker, kubernetes, raspberryPI*





influxdb

Different SKUs

- *Cloud*
 - Hosted by influxdata
- *OSS 2.X*
 - Current supported version
 - More powerful
- *OSS 1.X*
 - Currently deprecated
 - **Easier to use for relational people**



InfluxDB design principles

Time series data:

Series are stored in **measurements**

Measurements can have **fields** and **tags**

Series are sets of values of a single field, measured over time

Key of a series = [measurement name, time, tags, field name]

Example: CPU Usage

CPUHistory ← measurement name

Time

ServerName ← tag

CPUPercent ← field



Series – An Example

SQL Server CPU usage percent from `sys.dm_os_ring_buffers`

time	SQL Server CPU	Other CPU
08:00	10	3
08:01	15	2
08:02	18	6
08:03	73	1

In a relational database this
can be saved to a table: →

```
CREATE TABLE CPUHistory (  
    [time] time PRIMARY KEY,  
    SQLServerCPU tinyint,  
    OtherCPU tinyint  
)
```



Series – An Example

In InfluxDB this is saved in a **measurement** with **two series**:

time	SQL Server CPU	Other CPU
08:00	10	3
08:01	15	2
08:02	18	6
08:03	73	1

[CPUHistory, SQL Server CPU]

time	SQL Server CPU
08:00	10
08:01	15
08:02	18
08:03	73

[CPUHistory, Other CPU]

time	Other CPU
08:00	3
08:01	2
08:02	6
08:03	1



Series – Tags

If we add **tags**, we will get **one series for each tag value**:

time	ServerName	SQL Server CPU	Other CPU
08:00	ACCOUNTING	10	3
08:00	CRM	37	7
08:01	ACCOUNTING	15	2
08:01	CRM	48	5
08:02	ACCOUNTING	18	6
08:02	CRM	41	9
08:03	ACCOUNTING	73	1
08:03	CRM	28	7

[CPUHistory, **CRM**, OtherCPU]

time	Other CPU
08:00	7
08:01	5
08:02	9
08:03	7

[CPUHistory, **ACCOUNTING**, SQLServerCPU]

time	SQL Server CPU
08:00	10
08:01	15
08:02	18
08:03	73

[CPUHistory, **ACCOUNTING**, OtherCPU]

time	Other CPU
08:00	3
08:01	2
08:02	6
08:03	1

[CPUHistory, **CRM**, SQLServerCPU]

time	SQL Server CPU
08:00	37
08:01	48
08:02	41
08:03	28



InfluxDB – Storage Engine

- Series are stored independently
- We get one series for each combination of tag values → **cardinality**
- Does this look familiar? [Hint: columnstore indexes in SQL Server]
- Storage engine called **TSM** (Time series Structured Merge tree)
 - derived from **LSM** (Log-Structured Merge tree) → **Cassandra**
 - stores data in columns (series)
 - groups writes in memory and writes to disk in append mode
 - data on disk is always appended, never updated
 - series are compressed at regular intervals



Working with InfluxDB - Installation

- Download the binaries
- Expand the archive
- Check configuration
 - Port
 - Log file location
 - Authentication
- Run the daemon

DEMO!



Databases, retention policies, buckets...

SQL Server

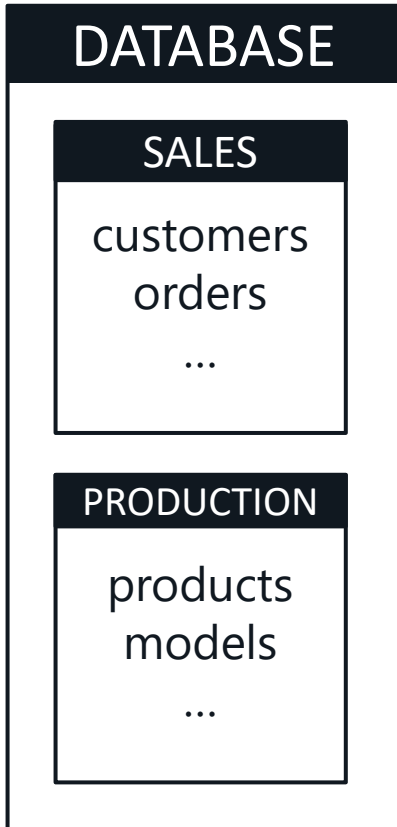
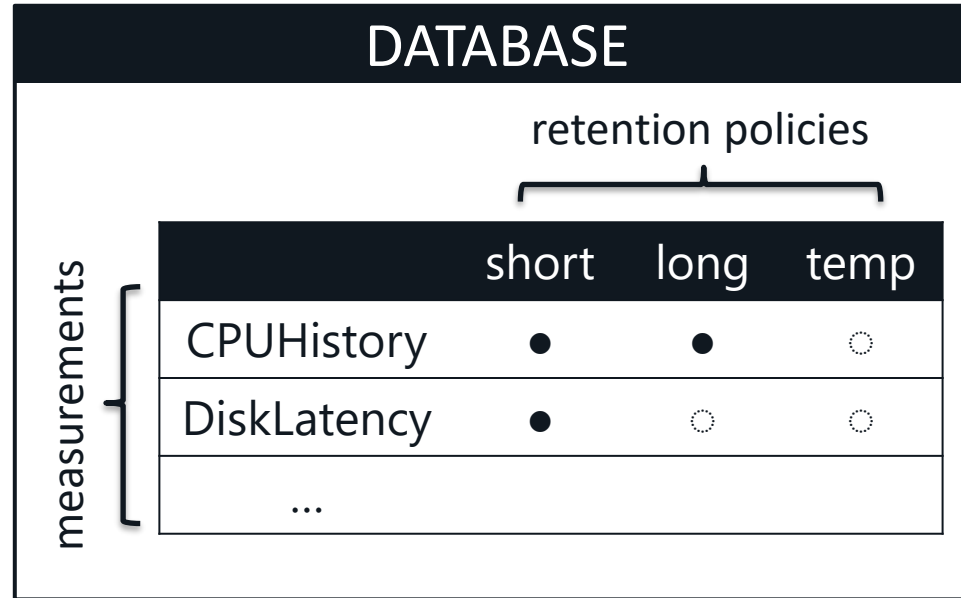


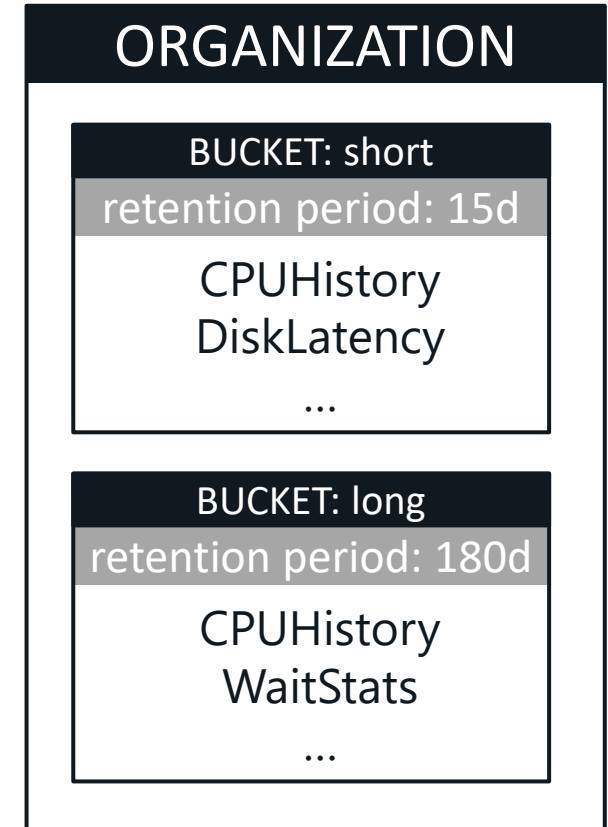
Table belongs to a single schema

InfluxDB 1.X



Measurement exists on all retention policies, you can write to whichever you need

InfluxDB 2.X



Measurement belongs to a single bucket



Working with InfluxDB – http API

InfluxDB can be queried and controlled with **http endpoints**

Endpoints are documented on the [reference page](#)

1.X and 2.X have different sets of APIs

1.X	2.X
v1 API	v2 API
v2 compatibility API	v1 compatibility API

Endpoints:	ping	Check status and version
	query	Query using InfluxQL or Flux (depending on version)
		Manage users, databases, ...
	write	Write data



Working with InfluxDB – CLI

InfluxDB ships with a command-line client, **influx.exe**

It is nothing but a client for the http API

Accepts several parameters:

<code>-host</code>	<code>-execute 'command'</code>
<code>-port</code>	<code>-type 'influxql flux'</code>
<code>-database</code>	<code>-format 'json csv column'</code>
<code>-username</code>	<code>-import</code>
<code>-password</code>	<code>...</code>

DEMO!



Query tools

There's more than the CLI: several query tools available:

- **Chronograf**

Official query and dashboard tool by Influxdata. Works in the browser

<https://www.influxdata.com/time-series-platform/chronograf/>

- **Time Series Admin**

Open Source query tool

<https://timeseriesadmin.github.io/>



Working with InfluxDB – Writing data

- Two ways to write to the database:
 - `/write` endpoint
 - using the CLI and the INSERT command
- Both expect the data to insert in a particular format called **line protocol**

DEMO!

`CPUHistory,server_name=ACCOUNTING SQLServerCPU=31,OtherCPU=7 1652386888000000000000`

The diagram shows the line protocol string `CPUHistory,server_name=ACCOUNTING SQLServerCPU=31,OtherCPU=7 1652386888000000000000` with four orange horizontal bars underneath it. Arrows point from labels below to these bars: **measurement** points to the first bar (`CPUHistory`), **tags** points to the second bar (`,server_name=ACCOUNTING`), **fields** points to the third bar (`SQLServerCPU=31,OtherCPU=7`), and **timestamp** points to the fourth bar (`1652386888000000000000`). Below the **tags** and **fields** labels is the text `n1=v1,n2=v2,...`.

measurement

tags
`n1=v1,n2=v2,...`

fields
`n1=v1,n2=v2,...`

timestamp
UNIX epoch ns in UTC



Working with InfluxDB – Writing data

Probably you don't want to write individual points → batches

```
CPUHistory,server_name=ACCOUNTING SQLServerCPU=31,OtherCPU=7 1652386888000000000000
CPUHistory,server_name=ACCOUNTING SQLServerCPU=48,OtherCPU=9 1652386903000000000000
CPUHistory,server_name=ACCOUNTING SQLServerCPU=53,OtherCPU=7 1652386918000000000000
CPUHistory,server_name=ACCOUNTING SQLServerCPU=47,OtherCPU=3 1652386933000000000000
...
```

Data types (for fields and tags):

numeric, string, integer, boolean

Timestamp is always in epoch format, from second to nanosecond

DEMO!



Working with InfluxDB – Writing data with telegraf

Telegraf is a data collection agent

Works with plugins :

INPUT:

Apache
Cassandra
Docker
Elasticsearch
Kubernetes
Nginx
PostgreSQL
SQL Server
... and more!

OUTPUT:

InfluxDB
Elasticsearch
File
Graphite
Graylog
Kafka
MongoDB
SQL
... and more!

DEMO!



Querying data with InfluxQL

InfluxDB supports 2 query languages:

InfluxQL → similar to SQL, born with version 1, still available in v2

```
SELECT MEAN(SQLServerCPU) AS avg_cpu  
FROM CPUHistory  
WHERE time > now() -1d  
GROUP BY server_name
```



Querying data with Flux

Flux → doesn't look like SQL, default in v2, also available in v1

Supports «advanced» query features like joins...

```
from(bucket: "demo/autogen" )  
  |> range(start: -1d)  
  |> filter(fn: (r) => r._measurement == "CPUHistory")  
  |> filter(fn: (r) => r._field == "SQLServerCPU")  
  |> mean()  
  |> group(columns: ["server_name"])  
  |> yield(name: "avg_cpu")
```



InfluxQL: SELECT syntax

SELECT <fields>

FROM <retention_policy>.<measurement>

[INTO <retention_policy>.<measurement>]

[WHERE <search_condition>]

[GROUP BY <tags> [FILL(<fill_expression>)]]

[ORDER BY <order_by_expression>]

[LIMIT <number>]

[OFFSET <number>]

[SLIMIT <number>]

[SOFFSET <number>]

[TZ(<timezone_clause>)]

DEMO!



InfluxQL: what about INSERT, UPDATE and DELETE?

INSERT does not exist

The write endpoint takes care of inserts

UPDATE does not exist.

To update a measurement, upload the data again.

Points that match the key [time, tags] will be updated


DELETE works pretty much like SQL

DEMO!



Retention policies and downsampling

- Default **retention policy** $\rightarrow \infty$
- Multiple retention policies, different retention periods
- **Continuous queries** aggregate and downsample the data to long term retention policy



Putting it all together

- Upload metrics using **Telegraf**
- Create a database in **InfluxDB**
- Display the data using **Grafana**



TIG stack

DEMO!



- Open source analytics/dashboarding platform
- Native InfluxDB support
- Powerful visualizations
- Easy to use



Conclusion

- Time series databases are great for storing and analyzing time series data
- Similarities with relational databases, but they are different beasts
- **InfluxDB** is the market leader and offers a great product with:
 - Optimized storage
 - Retention policies
 - Time series query capabilities
 - Continuous queries
- Great tool for IOT and telemetry observability solutions



Conclusion

- Time series databases can help you store, process and analyze huge volumes of time series data

Example:

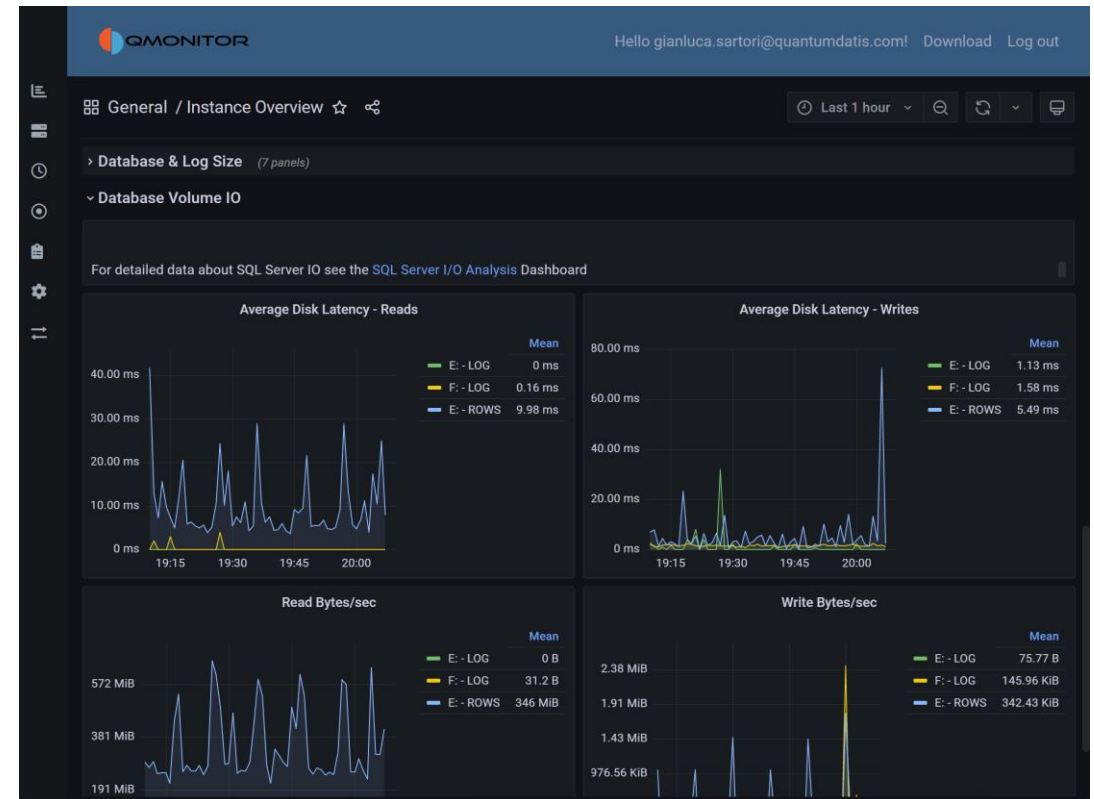


It's a complete SQL Server monitoring solution built on InfluxDB and Grafana

<http://qmonitor.quantumdatis.com>

User: demo@quantumdatis.com

Pass: qmonitor



Thank you!

Questions? Ask me!

spaghettidba@sqlconsulting.it

