

Time Series with SQL Server 2022

 DATA
SATURDAYS



Sponsors



With the support of:



Gianluca Sartori

Founder at Quantumdatis

Data Platform MVP

Works with SQL Server since version 7

DBA @ Scuderia Ferrari

Blog: spaghettidba.com

Twitter: [@spaghettidba](https://twitter.com/spaghettidba)



1. What is Time Series data?
2. Options on the market: InfluxDB
3. Storing Time Series data
4. Retention Policies & Downsampling
5. Querying Times Series data

What is Time Series data?

DATA
SATURDAYS



Definitions

What is **time series data**?

“Time series data (or time-stamped data) is a collection of observations for a single object (entity) at different time intervals.”

What is a **time series database**?

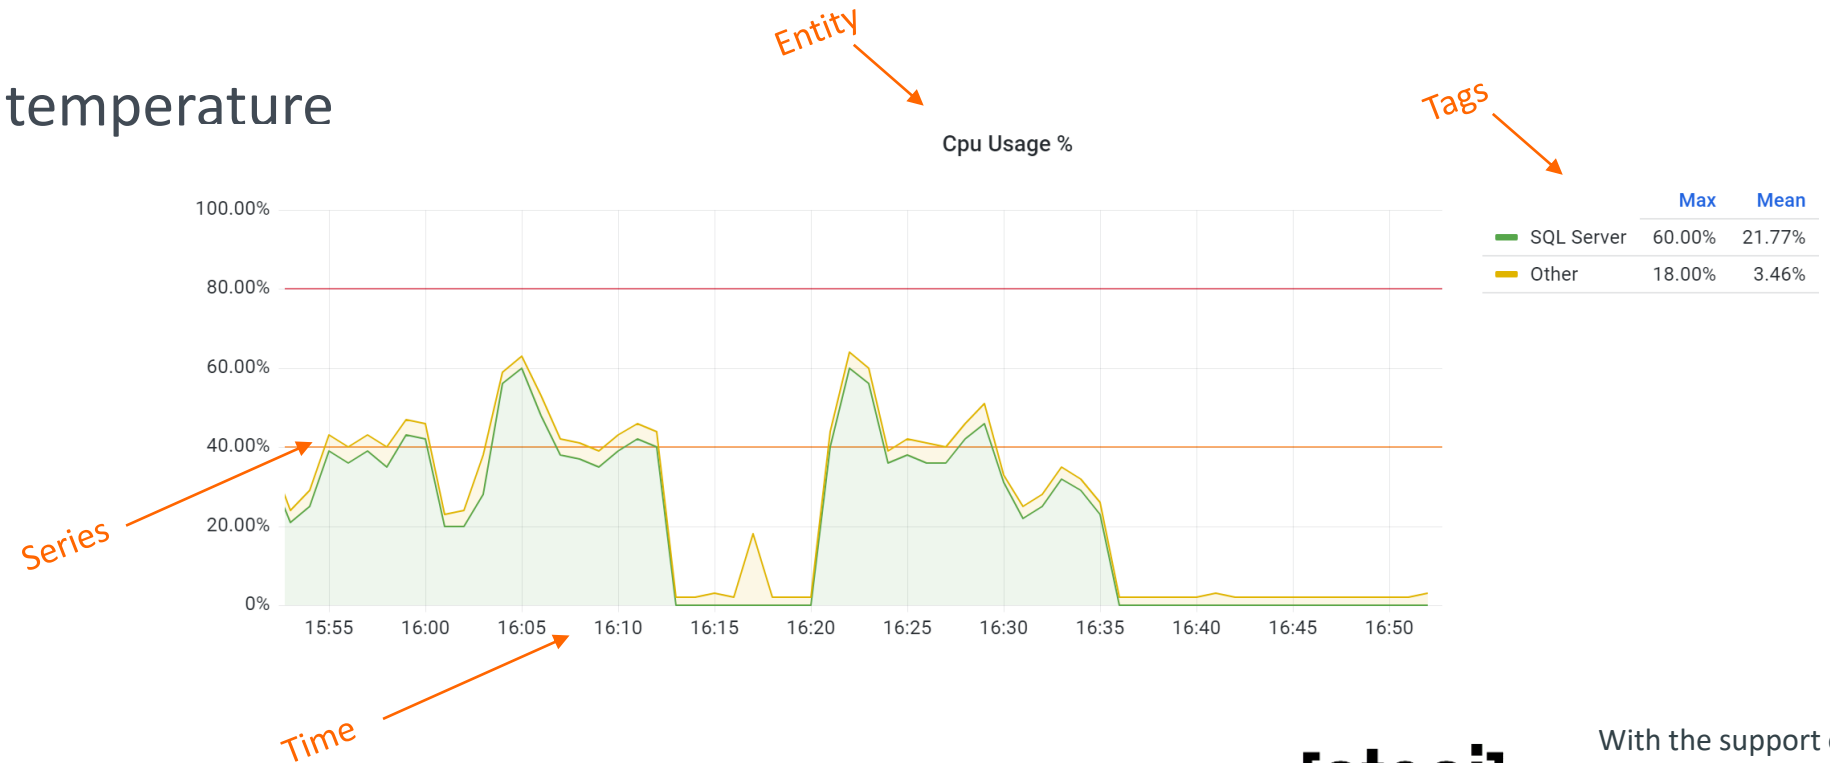
*“A time series database (TSDB) is a database optimized for **time series data** and for measuring change over time.”*

Time Series Data

Time series data is obtained by performing repeated measurements over time

Examples:

- Atmospheric temperature
- Stock prices
- CPU usage %
- Emails/sec
- Sensor data



Time Series Data


Metrics (Regular)

Measurements gathered at regular time intervals



Events (Irregular)

Measurements gathered at irregular time intervals



With the support of:

Time Series Data

How is this different from relational data?

- Continuous Stream of data in time order
- Bulk uploads of large sets of data
- High volume data
- Data is append-only – no updates
- Delete large volumes of data when it goes out of scope
- Downsampling and aggregating high resolution data to save space

Should I use a relational db, like SQL Server?

Relational databases do not deal very well with time series data

- No built-in storage optimizations
- No easy retention policies enforcement
- No time series specific query capabilities

- **Azure SQL Edge** is designed for time series data

See <https://docs.microsoft.com/en-us/azure/azure-sql-edge/>

- **SQL Server 2022** gets some of that: is it any good?

Options on the market: InfluxDB

DATA
SATURDAYS



Time Series Database Options

RANK	DBMS	SCORE		
		APR 2022	24 MOS ▲	12 MOS ▲
1	InfluxDB	30.02	+9.10	+2.85
2	Kdb+	8.78	+3.41	+0.52
3	Prometheus	6.31	+2.00	+0.55
4	Graphite	5.36	+1.90	+0.80
5	TimescaleDB	4.56	+2.79	+1.66
6	ApacheDruid	3.18	+1.27	+0.51
7	RRDtool	2.58	+0.06	+0.12
8	OpenTSDB	1.82	0.00	+0.02
9	DolphinDB	1.62	+1.23	+0.72
10	Fauna	1.42	+0.47	-0.07

Source: DB-Engines

39 Systems in Ranking, April 2022

InfluxDB



Most popular time series database

Designed and optimized for time series data

Uses optimized TSM storage engine with columnar compression

Built-in retention policies

Open Source with permissive licensing (MIT)

Written in GO

Cross platform: *windows, linux, macOS, docker, kubernetes, raspberryPI*

With the support of:

InfluxDB

Different SKUs

- *Cloud*
 - Hosted by influxdata
- *OSS 2.X*
 - Current supported version
 - More powerful
- *OSS 1.X*
 - Currently deprecated
 - **Easier to use for relational people**





InfluxDB design principles

Time series data:

Series are stored in **measurements**

Measurements can have **fields** and **tags**

Series are sets of values of a single field, measured over time

Key of a series = [measurement name, time, tags, field name]

Example: CPU Usage

CPUHistory ← measurement name

Time

ServerName ← tag

CPUPercent ← field

With the support of:



Series – An Example

SQL Server CPU usage percent from `sys.dm_os_ring_buffers`

time	SQL Server CPU	Other CPU
08:00	10	3
08:01	15	2
08:02	18	6
08:03	73	1

Storing Time Series Data

DATA
SATURDAYS






Series – An Example

In InfluxDB this is saved in a **measurement** (let's call it **CPUHistory**) with **two series**:

time	SQL Server CPU	Other CPU
08:00	10	3
08:01	15	2
08:02	18	6
08:03	73	1



With the support of:

Series – An Example



If we add **tags**, we will get **one series for each tag value**:

time	ServerName	SQL Server CPU	Other CPU
08:00	ACCOUNTING	10	3
08:00	CRM	37	7
08:01	ACCOUNTING	15	2
08:01	CRM	48	5
08:02	ACCOUNTING	18	6
08:02	CRM	41	9
08:03	ACCOUNTING	73	1
08:03	CRM	28	7

[CPUHistory, **ACCOUNTING**, SQLServerCPU]

time	SQL Server CPU
08:00	10
08:01	15
08:02	18
08:03	73

[CPUHistory, **ACCOUNTING**, OtherCPU]

time	Other CPU
08:00	3
08:01	2
08:02	6
08:03	1

[CPUHistory, **CRM**, OtherCPU]

time	Other CPU
08:00	7
08:01	5
08:02	9
08:03	7

[CPUHistory, **CRM**, SQLServerCPU]

time	SQL Server CPU
08:00	37
08:01	48
08:02	41
08:03	28



InfluxDB – Storage Engine

- Series are stored independently
- We get one series for each combination of tag values → **cardinality**
- Does this look familiar? [Hint: columnstore indexes in SQL Server]
- Storage engine called **TSM** (Time series Structured Merge tree)
 - derived from **LSM** (Log-Structured Merge tree) → **Cassandra**
 - stores data in columns (series)
 - groups writes in memory and writes to disk in append mode
 - data on disk is always appended, never updated
 - series are compressed at regular intervals

With the support of:

SQL Server 2022 – Storage Engine

Big Heaps or Clustered Indexes are problematic

- Huge in size
- Hard to query
- Hard to maintain

Alternatives?

- Data compression
- Clustered Columnstore Indexes


SQL Server 2022 – Columnstore Indexes

- Each column is stored independently
- Compresses data very efficiently
- Suitable when lots of repeating data exists
- Good for bulk operations
- Don't play well with UPDATES and DELETES

DEMO!

SQL Server 2022 – Columnstore Indexes

Data size for SQL Server performance counters, 2 instances, 15 days of data, 82 million rows



With the support of:


Retention Policies & Downsampling

DATA
SATURDAYS



Retention Policies & downsampling influxdb

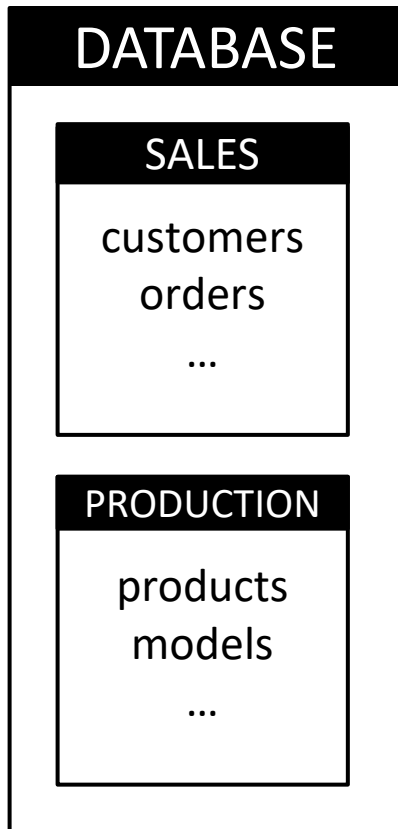
- Default **retention policy** $\rightarrow \infty$
- Multiple retention policies, different retention periods
- **Continuous queries**
aggregate and
downsample the data
to long term retention
policy



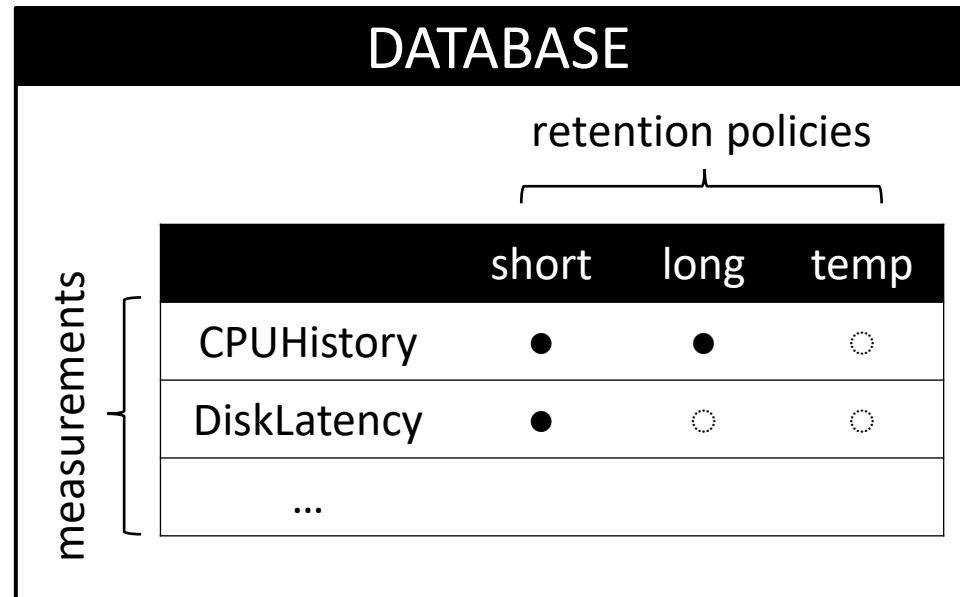
With the support of:

Databases, retention policies, buckets...

SQL Server

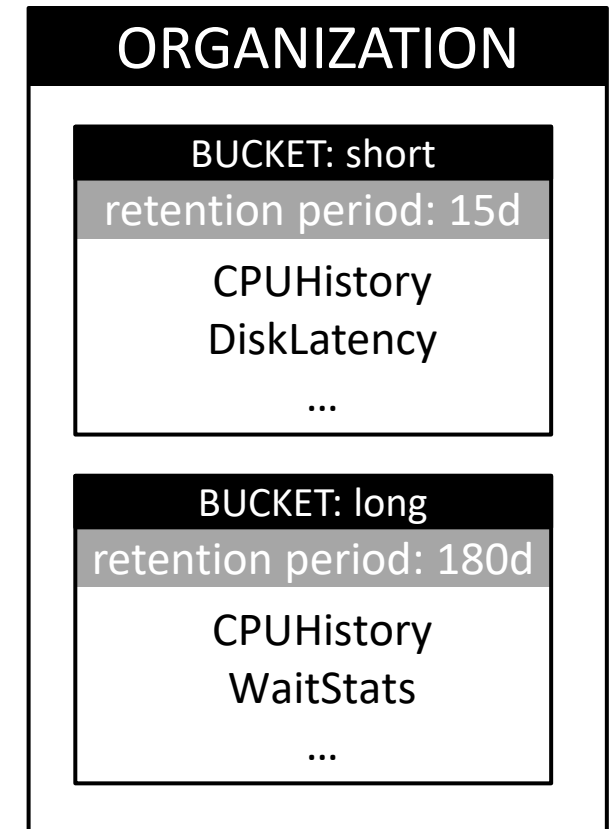


InfluxDB 1.X



Measurement exists on all retention policies,
you can write to whichever you need

InfluxDB 2.X



Measurement belongs to a single bucket
With the support of:

Retention Policies in Azure SQL Edge

Azure SQL Edge is a specialized SQL Server Edition for IoT applications
Designed with telemetry and Time Series data in mind
Collects data at the edge, uploads to the cloud
Supports DATA_DELETION:

```
ALTER TABLE [dbo].[sqlserver_performance]
SET (
    DATA_DELETION = ON (
        FILTER_COLUMN = [time],
        RETENTION_PERIOD = 1 month
    )
)
```

Retention Policies in SQL Server 2022

SQL Server 2022 gets new Time Series features...


```
ALTER TABLE [dbo].[sqlserver_performance]
SET (
    DATA_DELETION = ON (
        FILTER_COLUMN = [time],
        RETENTION_PERIOD = 1 month
    )
)
```

... does not get
this one though ☹️

```
Msg 102, Level 15, State 1, Line 117
Incorrect syntax near 'DATA_DELETION'.
```

Retention Policies in SQL Server 2022

- SQL Agent Job → DELETE
- Lots of rows → DELETE takes a long time, uses a lot of transaction log
- Columnstore Indexes don't like DELETES
- Partitioning allows switching data in and out with metadata-only operations
- Different partitions can have different compression settings
- Reduces efficiency of columnstore compression (each partition is independent)



DEMO!

With the support of:

Querying Time Series data

DATA
SATURDAYS



Querying data in InfluxDB



InfluxDB supports 2 query languages:

InfluxQL → similar to SQL, born with version 1, still available in v2

```
SELECT MEAN(SQLServerCPU) AS avg_cpu  
FROM CPUHistory  
WHERE time > now() -1d  
GROUP BY server_name
```

Querying data with Flux



Flux → doesn't look like SQL (more like kusto for Azure Data Explorer)

Default in v2, also available in v1

Supports «advanced» query features like joins...

```
from(bucket: "demo/autogen" )
  |> range(start: -1d)
  |> filter(fn: (r) => r._measurement == "CPUHistory")
  |> filter(fn: (r) => r._field == "SQLServerCPU")
  |> mean()
  |> group(columns: ["server_name"])
  |> yield(name: "avg_cpu")
```

With the support of:

InfluxQL: SELECT syntax



```
SELECT <fields>
  FROM <retention_policy>.<measurement>
  [ INTO <retention_policy>.<measurement> ]
  [ WHERE <search_condition> ]
  [ GROUP BY <tags> [ FILL(<fill_expression>) ] ]
  [ ORDER BY <order_by_expression> ]
  [ LIMIT <number> ]
  [ OFFSET <number> ]
  [ SLIMIT <number> ]
  [ SOFFSET <number> ]
  [ TZ(<timezone_clause>) ]
```

With the support of:



GROUP BY TIME

InfluxQL allows grouping by time:

```
SELECT MEAN(value) AS avg_lazy_writes_sec
FROM "sqlserver_performance"
WHERE counter = 'Lazy writes/sec'
      AND sql_instance = 'SQLCSRV04:SQL2017'
      AND time > now() - 1d
GROUP BY time(1h)
```

GROUP BY TIME



... GROUP BY time(1h)

time	lazy_writes_sec		time	avg_lazy_writes_sec
16-10-2022 10:00:00	398091	10:00	16-10-2022 10:00:00	398,091.00
...	...		16-10-2022 11:00:00	400,187.50
16-10-2022 10:59:59	395644	11:00	16-10-2022 12:00:00	406,095.00
16-10-2022 11:00:00	432211		16-10-2022 13:00:00	411,167.76
...	...	12:00	16-10-2022 14:00:00	417,963.40
16-10-2022 11:59:59	452325		16-10-2022 15:00:00	425,683.27
16-10-2022 12:00:00	456541	13:00	16-10-2022 16:00:00	433,001.40
...	...		16-10-2022 17:00:00	441,937.02
16-10-2022 12:59:59	433505	...	16-10-2022 18:00:00	449,937.73
16-10-2022 13:00:00	456654		16-10-2022 19:00:00	450,737.81
...	...		16-10-2022 20:00:00	453,467.40
16-10-2022 13:59:59	456788		16-10-2022 21:00:00	455,703.76
16-10-2022 14:00:00	417545		16-10-2022 22:00:00	457,270.19
...	...		16-10-2022 23:00:00	458,743.47
			17-10-2022 00:00:00	461,533.64
		

With the support of:

Introducing DATE_BUCKET

New DATE_BUCKET function in SQL Server 2022!!

```
SELECT DATE_BUCKET(hour, 1, time) AS time,  
       AVG(value) AS avg_lazy_writes_sec  
FROM [dbo].[sqlserver_performance_cci_partitioned]  
WHERE time >= '2022-10-01'  
       AND time < '2022-10-02'  
       AND counter = 'Lazy writes/sec'  
       AND sql_instance = 'SQLCSRV04:SQL2017'  
GROUP BY DATE_BUCKET(hour, 1, time)
```

DEMO!



FILL

InfluxQL allows filling intervals with data:

```
SELECT MEAN(value) AS avg_lazy_writes_sec
FROM "sqlserver_performance"
WHERE counter = 'Lazy writes/sec'
      AND sql_instance = 'SQLCSRV04:SQL2017'
      AND time > now() - 1d
GROUP BY time(1h) FILL(NULL)
```

FILL



... FILL(NULL | VALUE)

time	value
16-10-2022 10:35:45	398091
16-10-2022 10:36:01	398091
16-10-2022 12:16:15	406095
16-10-2022 12:16:30	406095
16-10-2022 13:06:45	411167
16-10-2022 13:07:01	411167

No data
for
11:00



time	avg_lazy_writes_sec
16-10-2022 10:00:00	398,091.00
16-10-2022 11:00:00	NULL
16-10-2022 12:00:00	406,095.00
16-10-2022 13:00:00	411,167.76

Introducing GENERATE_SERIES

New GENERATE_SERIES function in SQL Server 2022!!

```
SELECT * FROM GENERATE_SERIES(1, 100);
```

Generates a set of values from argument 1 to argument 2
Can be combined with DATE_BUCKET

DEMO!

	value
1	1
2	2
3	3
4	4
5	5
6	6
7	7
8	8
9	9
10	10
11	11

✓ Query executed successfully.

With the support of:

Conclusions

DATA
SATURDAYS



Is SQL Server 2022 good for time series data?

- Brings innovative T-SQL syntax to handle time series data
- Already has adequate storage capabilities
- Performance is acceptable
- Not as capable as a proper Time Series database like InfluxDB
 - Requires managing some things manually
(retention policies, downsampling)
- Offers T-SQL syntax and 360° support for relational data

THANK YOU!

Got questions? Email me!

Gianluca Sartori
Quantumdatis, owner

gianluca.sartori@quantumdatis.com
@spaghettidba

DATA
SATURDAYS

